

Wind: A Flow-based Programming Language

Joshua Weinstein

2018

Contents

0.1 Preface

This book is intended as an overview of the Wind programming language. It discusses the paradigm of flow-based programming, and the principles of it. The book discusses the advantages and disadvantages of flow-based programming, while then proceeding onto the syntax and usage of the Wind language. The book also, in detail, describes the C implementation of the language. The state management, the instructions, the translation of source code, and execution are all covered.

The Wind language is not a fully fledged, general purpose programming language. It is a language with a fundamental set of computation tools and components. Wind was developed with the following goals in mind.

Goals

1. An extremely light-weight language that is highly portable.
2. A programming language which does not use dynamic memory allocation.
3. A fluid, highly dynamically typed runtime.
4. A system that allows efficient transfers of immutable data.

The most unique element of Wind is that it is a "bare bones" language. It has no abstract syntax tree, no tokenizer, parser, or standard library. It reads and executes instructions directly from source code, which transition and alter several internal buffers. From a high level overview, these buffers pass around data to one another allowing the state of the data in the program to change. This choice of design permits the language to focus on manipulating and changing data, instead of managing and allocating resources.

Overall, the hope is for this book to serve as an inspiration or reference for computing and programming in a flow-based manner.

0.2 Dedication

This book is dedicated to my mother, Vida, who has guided me and raised me to be an incredible human being I am today. I would not be who I am today without her support.

0.3 About the Author

Joshua Weinstein Grew up in the sunny, green fields of Los Angeles, California. He went to the University of California Berkeley, where he studied Computer Science and Fine Art. He now works as a software engineer, specializing in distributed cloud systems and high performance search. In his spare time, he is an active contributor to the open source community.

Chapter 1

Flow-based Programming

Flow-based programming is a paradigm of programming that deals with the flow of data over one or more destinations. These destinations could in theory be any sort of structure. For this book, and the implementation of *Wind*, they are only considered buffers; fixed sized arrays controlled with pointers and indices.

This chapter will focus on an abstract perspective of flow-based programming. It will describe the fundamental concepts, such as nodes, data, and flows. It will also discuss the arrangements of how data flows can be constructed, and give some respective examples.

This chapter serves as an introduction to flow-based programming, and a preparation for the actual *Wind* language.

1.1 Data

In flow-based programming, data is considered any of the values or information that are computed or processed. Numbers, integers, strings, characters, lists, are all examples of data. Nearly all programming languages type the data that they contain and compute. At the raw level, data to a computer is simply sequences of 1's and 0's. Programming languages provide much more human readable forms of data. Integers or floats to hold numbers, strings to hold texts, and numerous data structures like a *list* to hold an ordered collection of data.

An important distinction in how data is typically assigned types is whether a type is mutable or immutable. An *immutable* type is one that cannot be altered or changed after it is created. In order to alter or change it, you must create a new copy of that immutable type from existing instances of the type. For example, an integer, or any number, is an immutable type. Numbers are created from operations performed on other numbers.

$$\begin{aligned}a &= 1 \\ b &= a + 1\end{aligned}$$

In the above example, a is a variable with the value of 1, while b is a variable with the value of the sum of a and 1. Neither of them can be changed. They might be able to be recreated under the same variable name, but this is not altering the existing value. Next, let's look at an example in C that deals with mutable and immutable data.

```
1 | int a = 1;
2 | int* b = malloc(sizeof(int));
3 | *b = a;
```

Here, b is an integer allocated dynamically on the program's heap, while a is a stack allocated variable. The main difference between them is that, b is a *mutable* data value that can be changed freely after it has been created. Yet a is immutable. The last statement, in the third, line sets the value of b to a . This means b can be considered a mutable type.

However, a and b are not both integers. b is a *pointer* to an integer, specifically a block of heap memory the size of one integer. In comparison to dynamically typed languages, b can be thought of as a one element list. In flow-based programming, the same data types have both mutable and immutable characteristics. Which characteristics of data determine its behavior is determined by the location and the *movement* of the data.

1.1.1 Data Movement

Generally, moving data in a program means the ownership of values and data are being passed around. Some languages define specific criteria for the movement of objects and memory, like C++. *Move semantics* usually refers to the process of moving dynamically allocated objects from one owner to another. Here is a simple example.

```
1 | int* a = malloc(sizeof(int) * 2);
2 | a[0] = 1;
3 | a[1] = 2;
4 | int* b = a;
```

The memory containing enough space for two integers, pointed to by a , is passed to b . After line four, both a and b point to or "hold" the same value. Accessing through either one yields the same behavior. Changing one of the integers of a changes it in b as well.

```
| printf("a is %d %d\n", a[0], a[1]);
| printf("b is %d %d\n", b[0], b[1]);
| b[0] = 5;
| printf("a is %d %d\n", a[0], a[1]);
```

Compiling this code will display the following in the terminal:

```
| 1 2
| 1 2
| 5 2
```

Moving the resource always involves sharing or copying a *pointer* or *reference* to some value. Movement never writes the value itself again in a new location. Thus, immutable data cannot be moved in the context of usual move semantics.

In flow-based programming, data is copied from one node or location to another. A reference to data is never moved from owner to owner. During that copy, the data is most often transformed in some way, such that the data written to a new location is not the same as the data in the old location. The following example illustrates how a flow-based program would transform and copy data.

Example: Two node transport

```

1 |
2 | void enterNode(int* node, int arg) {
3 |     *node = arg;
4 | }
5 |
6 | void transferToNode(int* startNode, int* endNode, int modify) {
7 |     *endNode = *startNode + modify;
8 |     *startNode = 0;
9 | }
```

Two functions above are shown, *enterNode* and *transferToNode*. Both of these functions demonstrate the acquisition and transfer of data between two nodes. It's important to note this is a very basic example. The nodes themselves here do not contain any information of how to "modify" data that they acquire and transfer. The *modify* parameter is simply used as an example of how the next node a data value is transferred to modifies each data point it receives. Take note again that no objects, pointers, or references enter the nodes, only literal data values.

1.2 Nodes and Locations

Nodes, destinations, and locations determine the path data will travel in a flow-based program. Nodes which receive, transform, and send data can take on an incredible amount of different forms. The form of a node is usually determined by the functionality or role it plays. While the structure of a node is quite variable, it must always abide by the following rules:

Rules of Flow Nodes:

1. *A node must possess some amount of readable and writable memory.*
2. *A node must have a valid memory address that can refer to itself.*
3. *A node must live longer than the data it handles.*

The first rule is the least ambiguous. All nodes must be containers, and be able to some non-zero amount of data. Specifically, nodes must own either an array or a pointer to a block of memory. If a node has to transport data along some path, it must be able to hold data temporarily at certain steps. For example, an `enum` value would not be a valid node, as even though it evaluates to an integer, it is not a container, nor does it possess the ability to own a block of memory.

Rule two is quite simple, a node must exist or be allocated in addressable memory. A *register* is not addressable, therefore a register cannot be a valid node. Constants defined with preprocessor statements like `#define` do not have a memory address and therefore cannot be used as nodes. Compile-time constants and values can be used as nodes if they are permitted for modification during runtime.

The last rule ensures that nodes will be able to support, transform and move data from the start of the program to the end of it. The key idea behind flow-based programming is that, nodes are the supporting component of data flow. Nodes need live and be reusable throughout data flowing through the program. Specifically, nodes should not be created dynamically as data is intended to be transported. When data enters a flow, all the nodes should already be initialized and ready to pass along data.

1.2.1 Node Memory

As explained earlier, nodes must possess or own modifiable memory. Nodes can store the incoming data they acquire in a variety of ways. An important factor to consider in how a node uses memory is what kind of data the node deals with. Further, if the data is all the same type, or if the node accepts multiple data types in the same pass.

First, let's look at a node which only holds a single value, and how it can be initialized, and deinitialized:

```

1 | #include <string.h>
2 |
3 | typedef struct
4 | {
5 |     void* data;
6 |     size_t size;
7 |     int occupied;
8 | } SingleNode;
9 |
10 | void SingleNode_init(SingleNode* node, size_t size) {
11 |     node->size = size;
12 |     node->data = malloc(size);
13 |     node->occupied = 0;
14 | }
15 |
16 | void SingleNode_deinit(SingleNode* node) {
17 |     free(node->data);
18 |     memset(node, 0, sizeof(SingleNode));
19 | }
```

Let's break down the above code. First, we see the `struct` describing the attributes of the node. It has a `void` pointer to a block of memory, a field indicating the size, and a field acting as boolean to check if the node currently possess a data value. This implementation allows any C type to be placed and written into the memory of the node. However, in order to access the data we

must cast the data pointer as some C type. While this isn't problematic, it urges the need for the types of data to be tracked externally to the node.

In this setup, initializing a node requires that the node itself, the scalar type, already be declared and stored somewhere. The pointer to the node body is passed in, initializing it with its memory and size. Then, once the node is done being used in the program, or needs to be resized, it can be de-initialized, by freeing the memory and clearing the data in the node's body. We want to permit the node itself to be stored in a centralized location, rather than allocating the node itself with a call to `malloc()` each time it is initialized.

A downside to this approach is that it decouples what is essentially the metadata of the node, and the data of the node. It requires two steps to have a fully functional node. Next, let's look at nodes that combine both its data and information in a single scalar body.

```

1 | #include <stdlib.h>
2 |
3 | typedef struct
4 | {
5 |     size_t capacity;
6 |     size_t len;
7 |     unsigned char[];
8 | } BlockNode;
9 |
10 | BlockNode* BlockNode_new(size_t size) {
11 |     BlockNode* node = malloc(sizeof(BlockNode) + size);
12 |     node->capacity = size;
13 |     node->len = 0;
14 |     return node;
15 | }
```

Note: The above uses a flexible array parameter, a special array available since C99. Flexible array parameters must always be the last members of a struct, and cannot be directly used with the `sizeof` operator.

BlockNode is a node that uses a dummy pointer to refer to a variably sized body of memory it carries at the end of itself.